

JING-ZHANG CIVIC AI LOOP

A CIVIC PROVING GROUND FOR URBAN AI

First deliver a civic proof space anyone can traverse, stop, appeal and replay—then discuss eleven kilometres.

ONE LINE · THREE STATIONS · ONE RECEIPT



v2.9 COUNTER-EVIDENCE GATE | BLOCKED_BEFORE_G2 | 0/6 PASS - 6/6 BLOCK

Provisional boundary | zero-feature constraints | operator and data governance pending | no field baseline | P1-P3 not executed.
Desktop replay is not field authorisation. Clear every block before human G2 review.

EXECUTIVE SUMMARY

The First 100m is the only launch action; three stations and the full corridor must earn scale through independent replay.

01 SPATIAL THESIS

Railway station, line, timetable and handover become entry, takeover, appeal and exit for urban AI.

02 LAUNCH BOUNDARY

Only one accessible public task chain P1; P2 and P3 do not compete for launch resources.

03 FAILURE STRATEGY

If the ordinary path, staffed desk or appeal point fails, the run fails.

04 SCALE VERDICT

Independent review returns expand, repair or exit; a demo is not deployment permission.

EVIDENCE MATURITY

For every claim: what is its source, what can it support, and what evidence replaces it?

01 E3 MANDATE

Brief and policy support scope and principles.

02 E2 REPOSITORY

Site package and fact guide support research context.

03 E1 RECOMPUTABLE


Provisional geometry supports concept comparison and consistency only.

04 E0 TO SURVEY


Tenure, utilities, cost, operator and performance stay unknown.

THREE-SCALE WORKING FRAMEWORK

First 100m → three stations → full corridor: each scale has prerequisites and explicit unknowns.




ONE LINE · THREE STATIONS · ONE RECEIPT



THE FIRST 100M CIVIC PROOF SEGMENT

A design name for the first continuous proof unit—not an unverified measurement.




SPATIAL ACCEPTANCE LOGIC

Choose → assist → handoff → challenge → replay. If any step fails, the run fails and does not scale.

SPATIAL + LAND-USE STRUCTURE

Prioritise existing ground floors, railway nodes, under-bridge and leftover spaces—not demolition-led expansion.



WALKABILITY, BLUE-GREEN + REVERSIBLE TECH

After technology is removed, walking, rest, lighting, staffed service and railway narrative still work.

01 ORDINARY ROUTE

Equally visible and always available.

02 CROSS LINKS

Link neighbourhoods to the civic spine; avoid a tech island.

03 BLUE-GREEN BUFFER

Serve everyday comfort and safety before experiments.

04 ROBOT BOUNDARY

Controlled testing only: limited zone, time and human takeover.

ROLES OF THREE ZONES + TWO WINGS

Not five park labels, but an evidence relay: verify, translate, experience, service and test.

01 ZHONGZHIYUAN

Joint stack verification; asks whether it can work safely.

02 AI ORIGIN

Open translation and civic co-test; asks whether people can use and challenge it.

03 DAZHONGSI

Daily consumption and public experience; asks whether rights hold.

04 TWO WINGS

Professional services + bounded field scenes, without inventing partner commitments.

P1 PROTOCOL + GATES

Replace data, walk the site, establish non-AI baseline, staff service and rehearse synthetically before civic tasks.



METRICS: UNKNOWN IS A RESULT

Do not convert unmeasured into passed, hide the worst group with averages, or call a demo field performance.

Accessible continuity	UNKNOWN TO MEASURE
Human handoff rate	UNKNOWN TO MEASURE
Severe factual error rate	UNKNOWN TO MEASURE
Worst-group performance gap	UNKNOWN TO MEASURE
Near-miss and takeover events	UNKNOWN TO MEASURE

ACCESSIBILITY + WORST GROUP

Wheelchair users, older adults, carers, visitors and developers run the same task; scale only what the worst group can complete.

01 SAME TASK

No incomparable demo scripts for different groups.

02 SAME RIGHT TO EXIT

Refusing AI must not reduce access to public service.

03 SMALL-CELL PROTECTION

De-identified grouping only when lawful and necessary; suppress small cells.

04 VISIBLE FAILURE

Show blockers and remedy; never hide them in a composite score.

PHASING + REVERSIBLE DELIVERABLES

The first 90 days establish facts, baseline, human fallback and rehearsal only; quantities and cost await professional verification.

01 DAYS 0—30

Replace official data, verify rights, walk the site and assign service responsibility.

02 DAYS 31—60

Non-AI control, staff roster, accessible task script and synthetic data.

03 DAYS 61—90

Closed rehearsal, independent review and repair; no opening if failed.

04 EXIT RESOURCE

Reserve resources for removal, data handling, staffed service and public explanation.

RUN RECEIPT + ACCOUNTABILITY

Each run leaves minimal replayable, accountable and expiring evidence—without unnecessary identity.

VERSION + SOURCE

MINIMAL · REPLAYABLE · EXPIRING

TASK + OUTCOME

MINIMAL · REPLAYABLE · EXPIRING

HUMAN OWNER + HANDOFF

MINIMAL · REPLAYABLE · EXPIRING

WORST GROUP + FAILURE

MINIMAL · REPLAYABLE · EXPIRING

APPEAL + EXPIRY + DISPOSAL

MINIMAL · REPLAYABLE · EXPIRING

EXPAND / REPAIR / EXIT

MINIMAL · REPLAYABLE · EXPIRING

RISKS, STOP LINES + VERDICT

The Loop is not about adopting AI faster; it gives the city spatial and institutional capacity to reject unfit AI.

01 BOUNDARY RISK

Provisional geometry mistaken for legal boundary → label across every carrier and replace.

02 RIGHTS RISK

Ordinary route, human handoff or appeal blocked → stop immediately.

03 EVIDENCE RISK

Not replayable, unclear provenance or severe error → no scaling.

04 FINAL VERDICT

EXPAND | REPAIR | EXIT; there is no automatic deployment.